


④ The Abstraction: The Process



- * The process is just a running program, but also a fundamental abstraction from the OS.
- * The OS takes the bytes of a program & helps give them meaning by executing them.
- * Typical OS may be running tens/hundreds of processes @ once.
 - ↳ don't have to worry if a CPU is available.

Coux: How to provide the illusion of many CPU?

- ↳ do so by virtualizing the CPU, i.e., running one process, stopping it, and running another.
- ↳ this technique is known as time-sharing the CPU.
 - ↳ performance cost of timesharing is that processes will run more slowly
- * A mechanism is low-level methods/protocols that implement a needed piece of functionality.
 - ↳ necessary for implementation of CPU virtualization
 - ↳ A context-switch gives the OS the ability to stop running one program & start running another on a given CPU.

Tip: Space Sharing

- * Space sharing is when a resource, i.e. disk space, is divided (in space) among those who wish to use it.
- * OS implements policies (algorithms) for making decisions.
 - ↳ Scheduling Policy determines which program to run on a given CPU.

4.1 Abstraction: A Process

- * We can summarize a process by the different parts of the system it accesses/affects.
- * A process' machine state is what a program can read/update & what parts of the machine are important for its execution
 - ① Process' memory (address space) are the locations (addresses) in memory that the program reads/writes instructions are.
 - ② Registers are included as many instructions access them
 - ↳ program counter
 - ↳ stack & frame pointer (func. params, local vars, return addresses)

Tip: Separate Policy & Mechanism

- * Typical design paradigm of OS is to separate high-level policies from low-level mechanisms.
(which) (how)
 - ↳ allows for modularity — don't have to change mechanisms if want to change policies

4.2 Process API

- * following are descriptions of common functionality in a process API:
 - ① Create: OS must have some method to create a process
 - ② Destroy: Many processes terminate themselves when done running, but there should be a method for a user to kill them just in case
 - ③ Wait: Can be useful to wait for a process to stop running
 - ④ Miscellaneous:
 - ↳ suspend?
 - ⑤ Status:
 - ↳ runtime duration, state of process

4.3 Process Creation


* Q. How are programs transformed into processes?

① OS must load instructions & static data into memory (address space of process)

↳ OS must read bytes of an executable from disk/SSD and load to mem.

↳ Old OS load programs eagerly, i.e. all at once before executing.

↳ Modern OS load lazily, i.e. load instructions/data as needed during execution.



② Allocate memory for the program's stack

↳ local vars, function params, return address

↳ initialize the stack of arguments (fill in argv & argc)

③ Initialization related to I/O

↳ setup file descriptors?

④ Start execution @ entry point

↳ relinquish control of CPU to the running process.

4.4 Process States

* States that a process can be in @ any given time:

* Running: A process is running on a processor

* Ready: A process is ready to run but OS has chosen not to run it yet.

* Blocked: process performed some operation st. it is not ready to run until some other event takes place


↳ e.g. I/O request to a disk

↳ this means some other process can use the processor

* A process can be moved between $\text{Ready} \rightleftharpoons \text{Running}$ @ discretion of OS.

↳ Ready \rightarrow Running \Rightarrow Process was scheduled

↳ Running \rightarrow Ready \Rightarrow Process was descheduled



describes how OS transitions process states.

NOTE: when I/O done doesn't mean OS will immediately run the process that requested

4.5 Data Structures

* OS maintaining some kind of process list to track the state of each process.

↳ track ready processes & info for running processes.

↳ must have a way of tracking blocked processes to know which ones to wake when their event completes.

* Xv6 Kernel tracks:

* Register context: the state of the registers when a process becomes blocked.

↳ allows for a process to resume exactly where it left off

* pid, parent process, cwd, start & size of address space in mem, etc.

* process' state

↳ Zombie State: used by some OS to indicate a process has finished but info hasn't been cleaned up.

↳ maybe parent process doesn't need to wait for child so doesn't need to examine & clean its final state.

Note: Sometimes parent needs to wait for child process to examine its exit code, and will signal to OS to clean up its info after finished examining it!

Aside: The Process List

- * Something similar is necessary for any time-sharing OS to complete context switches.
- * Often refer to individual entries in the list as Process Control Blocks (PCB)

4.6 Summary

* will learn:

- ① low-level mechanisms to implement processes
 - ② High-level policies to schedule processes
- These will teach us how OS virtualizes CPU