




Introduction

KEY

-  : important
-  : definition
-  : comment

Operating Systems
(Three Easy Pieces)

Matthew Bourque 

PREFACE

* three major elements:

- ① Virtualization
- ② Concurrency
- ③ Persistence

* Each major section presents an abstraction that the text goes further to describe the mechanisms underneath.

* Most important practice are the projects

↳ [github: renzi - arpaudusseau/ostep-projects](#)

* Worth taking a look into the original sources where the ideas described in the book came from.

" LEARN BEYOND THE CLASSROOM! "

DEALING ON THE BOOK

* Through the three major elements, we will learn:

- * how OS decides what program to run next on CPU
- * " " handles memory overload in a virtual memory system
- * how virtual machine monitors work
- * how to manage info on disks
- * how to build a distributed system that performs despite parts having failed.

" I SEE AND I REMEMBER; I DO AND I UNDERSTAND "

INTRO TO OS

* Patt & Patel [CPO3]

* Bryant & O'Hallaron [BOH10]

* Running programs execute instructions:

- ① fetches instruction from memory
- ② decodes instruction
- ③ executes it.

} Von
Neumann
model

* The Operating System is a body of software that is responsible for making it easy to run programs

- ① run programs
- ② allow programs to share memory
- ③ allow programs to interact w/ devices.
- :

* OS does this through virtualization: the process of taking a physical resource & transforming it into a more general/powerful virtual form of itself.

* OS provides an interface (API) st. a user can make use of its features (system calls) which collectively form a standard library.

* Virtualization allows:

- ① Many programs to run
- ② " " " concurrently access their own instructions
- ③ " " " access devices

⇒ the OS is known to be a resource manager

* CPU, memory, disk are all resources OS uses

CPU: how does OS attain virtualization?
how does it do it efficiently & what
are the hardware requirements?

Virtualizing The CPU

CPU: Consider a system w/ a single processor, the OS can virtualize the CPU and make it seem as if there exists multiple processors when running multiple programs "simultaneously"

↳ Not really simultaneous, the OS has a scheduling mechanism which dictates which program should run & quickly switches between programs. (Scheduling Policy) CPU executes

Virtualizing Memory

* Modern machines model physical memory through an array of bytes.

Address	Data
0x00	32
⋮	⋮
0xFF	

- * Memory is accessed all the time when running a program
 - ↳ where instructions are loaded/stored
 - ↳ where data structures are kept
- * When we run two instances of mem.c "simultaneously" we can see:
 - ① Each instance has its own unique pid (process identifier)
 - ② Each process seemingly has its own private memory, instead of sharing the same memory w/ the other processes.
 - ↳ In reality, they are sharing the same physical memory (RAM)
- * OS is virtualizing memory as each process gets its own virtual address space which OS maps to the physical memory

Concurrency

- * concurrency is working on many things @ once, often in the same program.
 - ↳ concurrency brings forth issues that ought to be addressed
- * threads.c highlights a problem:
 - ↳ start by creating two threads: a function running in the same memory space as other functions, w/ more than one of them active @ a time.

Same
virtual
address
space

- * When executing `threads.c` w/ sufficiently large input, we obtain data races because incrementing an integer variable is not atomic, i.e., it is possible to switch between threads before the operation is done.

↳ incrementing will take three instructions:

- ① load data from var
- ② increment data
- ③ save data from var

} → NOT ATOMIC!

CPUX: What mechanisms can we use to ensure concurrently executing threads will behave in the way we expect.

[lock the OS from switching between threads until the dangerous operation is done?]

Persistence

- * In system memory, data can be lost
- * Some devices (DRAM) store data in unstable ways
 - ↳ loss of power ⇒ loss of data
- * Need for hardware/software that can store data persistently: over long periods of time & across power losses.
- * Hardware (I/O devices):
 - ① Hard Drive (long lived info)
 - ② Solid-State Drive (SSD)

} Persistent Storage
- * Software (File Systems):
 - ↳ reliably & efficiently store any user files.
 - ↳ OS assumes users want to store info in files.

* system calls request some action to be done from the OS.

↳ open() — opens & creates file

↳ write() — writes data to file

↳ close() — closes the file, i.e., no longer accessing it

↳ * These system calls are routed to the file system, i.e., the part of OS that will handle these requests

* To do so, FS must:

+ figure out where new data will reside on disk

↳ issue I/O request to storage device

* most FS delay writes to do them in batches (more efficient)

* " " have protocol to ensure system can recover some info if failure were to occur during write.

* encode itself in efficient data structure (advanced b-tree)

Design Goals

① build abstractions that make OS easier to use

② provide high performance [minimize overheads of OS: excess/indirect
computation time / space usage]

③ provide protection between applications

↳ prevent malicious programs from harming other programs / OS itself while allowing multiple programs to run @ once.

↳ do so through isolation of processes

④ provide reliability and ensure program doesn't crash & lose user's data.

⑤ energy-efficient

⑥ mobility (iphones)

History

① Early OS : Just Libraries

- * essentially just a set of libs of common functions
- * usually on these systems only one program ran @ a time
↳ humans would decide what order to execute jobs in.
- * too expensive to let single user interact w/ machine so ran jobs in batches.

② Beyond Libs : Protection

- * realized that code ran on behalf of OS was special as has control over devices. (should be treated diff)
- * system call was designed to make transition to OS more controlled through hardware instruction (traps)

↳ transfer control over to OS while raising hardware privilege level

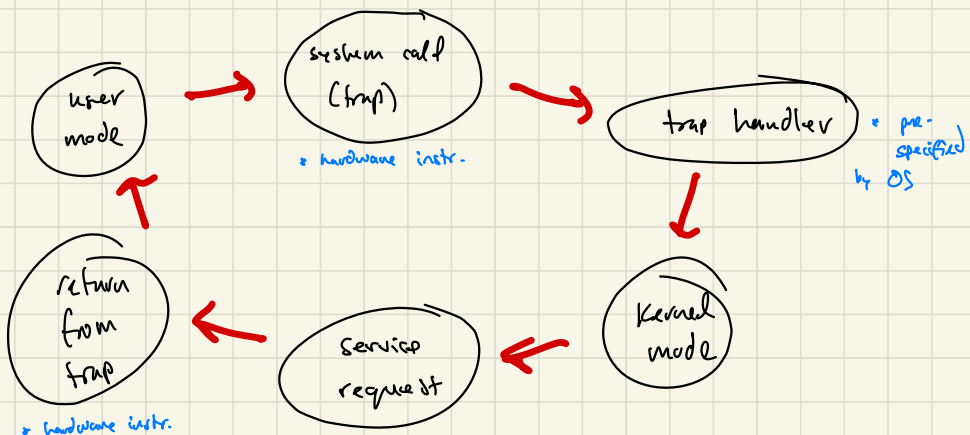
① kernel mode

② user mode (hardware restricts what applications can do)

↳ no I/O request to disk

↳ no access to physical memory

↳ can't send packet on network.



③ Multiprogramming

- * creation of minicomputer \Rightarrow more access to comp. \Rightarrow more need for OS
- * multiprogramming developed to better use machine resources.
 - \hookrightarrow load a bunch of jobs & switch rapidly btwn them.
 - \hookrightarrow switching important bc I/O requests are SLOW
 - \hookrightarrow instead of waiting, switch to another program on CPU
 - \hookrightarrow led to concurrency issues & memory protection issues

④ Modern Era

- * creation of PC \Rightarrow even more accessibility to comp.
- * microsoft DOS (Disk OS) didn't focus on memory protection.
- * Mac OS had cooperative scheduling \Rightarrow running thread force reboot
- * today OS are similar to minicomputer & focus on design goals.

Summary

* This book won't go deep into:

- ① Networking
- ② Graphics
- ③ Security

* Homeworks

- ① Simulations
- ② Real-world code